

ZAIN SHAFIQUE

+923047549668 | zainshafique23@gmail.com | [LinkedIn](#) | [Github](#)

Experience

Data Analyst | Remote, Women in Sport × Statistics Without Borders 09/2025 – 04/2026 | USA

- Analyzed survey data from ~4,400 youth respondents using Python, and SPSS, using cross-tabulations and visualizations to identify participation barriers by gender and ethnicity
- Co-authored a team report on safety, confidence, and health-related barriers (e.g., 45.7% of Asian females did not feel safe exercising outdoors), included as an appendix to the final report delivered to Women in Sport

Freelance Financial Operations Specialist | Remote, US Restaurant Franchise Client 01/2024 – Present | USA

- Manage accounting and financial reporting for 4 restaurant locations using QuickBooks, including bookkeeping, reconciliations, and insurance administration
- Prepare monthly P&L and sales reports for ownership across all locations

AI Engineer Intern | Remote, Verior 06/2025 – 08/2025 | Pakistan

- Built a Flask microservice using the Cohere API for text generation and summarization, with a modular LLM wrapper, deployed on Vercel
- Built a Node.js/Express translation API supporting 20 languages, with a JavaScript front end, deployed on Vercel

Data Science Intern | Remote, Bytewise Limited 06/2024 – 09/2024 | Pakistan

- Built ETL workflows for data preparation and implemented sentiment analysis and anomaly detection models in Python

Projects

ADF Pipeline Patterns — CDC, REST API & File Routing

- Built 7 production-grade ADF pipeline patterns including CDC incremental load (JSON & DB-based audit logging), dual REST API pagination, dynamic file routing, dynamic schema mapping, Delta Lake upserts via Mapping Data Flow, and multi-table orchestration driven by a SQL control table with parallel execution.

End-to-End Azure Data Lakehouse - ADF to Power BI

- Built a production-grade end-to-end data engineering pipeline on Azure using Adventure Works data, orchestrating the full Medallion Architecture (Bronze → Silver → Gold) across Azure Data Factory, Databricks, and Synapse Analytics Serverless SQL
- Implemented enterprise-level enhancements including Delta Lake storage in the Silver layer, parameterized Databricks notebooks with secret scope authentication, Great Expectations data quality checks, and CI/CD deployment via GitHub Actions
- Delivered analytics-ready data surfaced through a Power BI dashboard, with T-SQL views built on Synapse OPENROWSET for schema-on-read querying

Databricks Medallion Pipeline

- Engineered an end-to-end data lakehouse pipeline on Databricks (Bronze → Silver → Gold) over e-commerce sales data using PySpark and Delta Lake
- Implemented incremental watermark loading and modeled the Gold layer as a Kimball-style star schema with SCD Type 1 & 2 for full historical dimension tracking
- Leveraged Delta Lake ACID-compliant storage to ensure reliable, analytics-ready datasets for downstream reporting

Medallion Data Warehouse - SQL Server

- Built an end-to-end data warehouse on SQL Server implementing Medallion Architecture (Bronze → Silver → Gold), integrating ERP and CRM source data via T-SQL ETL pipelines
- Designed a star schema (fact & dimension tables) in the Gold layer with data cleansing and standardization in Silver, enabling optimized analytical queries

Education

National University of Sciences & Technology (NUST)

08/2019 – 06/2023 | Islamabad, Pakistan

Bachelor of Engineering in Software Engineering

Technical Skills

Languages & Libraries: Python (Pandas, NumPy, Scikit-learn, TensorFlow), SQL, T-SQL, PySpark

Cloud & Data Engineering: Azure Data Factory, Azure Databricks, Azure Synapse Analytics, ADLS Gen2, Azure SQL DB, Azure Logic Apps, Great Expectations, SQL Server, Git/GitHub

Data Modeling & BI: Medallion Architecture, Kimball Star Schema, SCD Type 1 & 2, ETL/ELT, CDC, Power BI, DAX, REST APIs